


# Linear Classification

**Kietikul Jearanaitanakij**

Department of Computer Engineering, KMITL

# Regression and Classification with Linear Models

- Regression analysis is a statistical process for estimating the relationships between a dependent variable  $y$  and independent variable(s)  $x$ .
- Given data points of  $(x, y)$ , linear regression consists of finding the best-fitting straight line (function) through the points. The best-fitting line is called a regression line.




# Univariate linear regression

- Univariate linear function (a straight line) with input  $x$  and output  $y$  has the form  $h_w(x) = w_1x + w_0$ , where  $w_0$  and  $w_1$  are real-value coefficients (weights) to be learned.
- To find  $h_w$  that best fit input data  $x$ , we have to find the values of the weights  $[w_0, w_1]$  that minimize the loss (e.g. training errors).

Number of examples  $\rightarrow$   $N$       Target value  $y_j$       Actual value  $h_w(x_j)$        $\text{mean}(\text{true} - \text{pred})^2$

$$\text{Loss}(h_w) = \sum_{j=1}^N \left( y_j - h_w(x_j) \right)^2 = \sum_{j=1}^N \left( y_j - \underbrace{(w_1x_j + w_0)}_{h_w(x_j)} \right)^2$$



Note that this loss function is convex, with a single global minimum.

- The sum  $\sum_{j=1}^N (y_j - (w_1 x_j + w_0))^2$  is minimized when its partial derivatives with respect to  $w_0$  and  $w_1$  are zero.

$$\frac{\partial}{\partial w_0} \sum_{j=1}^N (y_j - (w_1 x_j + w_0))^2 = 0, \quad \frac{\partial}{\partial w_1} \sum_{j=1}^N (y_j - (w_1 x_j + w_0))^2 = 0$$

Optimization problem: solve equations for  $w_0$  and  $w_1$

សម្រាប់ស្វែងរក  $w$  ត្រឹមត្រូវ

- For example in the above figure, the solution is  $w_1=0.232$ ,  $w_0= 246$ , and the line with those weights is shown as a dashed line in the figure.

- In a general optimization problem, this optimization problem can be addressed by a **hill-climbing algorithm** that follows the gradient of the function to be optimized. We will use **gradient descent** to minimize the loss. *Algo for derivative*

$w$  = any point in the parameter space

#  $w = \{w_0, w_1, w_2, w_3, \dots, w_n\}$

loop until convergence

for each  $w_i$  in  $w$  do

$$w_i \leftarrow w_i - \alpha \frac{\partial}{\partial w_i} \text{Loss}(\mathbf{w}) \quad (1)$$

*Note:*  $\alpha$  is usually called the learning rate.

- For univariate regression, the loss function is a quadratic function, so the partial derivative will be a linear function.

$$\begin{aligned}
 \frac{\partial}{\partial w_i} \text{Loss}(\mathbf{w}) &= \frac{\partial}{\partial w_i} (y - h_w(x))^2 \\
 &= 2(y - h_w(x)) \times \frac{\partial}{\partial w_i} (y - h_w(x)) \\
 &= 2(y - h_w(x)) \times \frac{\partial}{\partial w_i} (y - (w_1 x + w_0))
 \end{aligned}$$

- Applying this to both  $w_0$  and  $w_1$  we get:
- $$(1) \begin{cases} \frac{\partial}{\partial w_0} \text{Loss}(\mathbf{w}) = -2(y - h_w(x)) \\ \frac{\partial}{\partial w_1} \text{Loss}(\mathbf{w}) = -2(y - h_w(x)) \times x \end{cases}$$
- Handwritten notes and diagrams:
- For  $\frac{\partial}{\partial w_0}$ :  $2(y - h_w(x)) \times \frac{\partial}{\partial w_0} (y - (w_1 x + w_0))$ . The derivative of  $w_0$  is 1, so it becomes  $2(y - h_w(x)) \times (0 - (0 + 1)) = -2(y - h_w(x))$ .
  - For  $\frac{\partial}{\partial w_1}$ :  $2(y - h_w(x)) \times \frac{\partial}{\partial w_1} (y - (w_1 x + w_0))$ . The derivative of  $w_1$  is  $x$ , so it becomes  $2(y - h_w(x)) \times (0 - (x + 0)) = -2(y - h_w(x)) \times x$ .
  - The function  $h_w(x)$  is defined as  $h_w(x) = w_1 x + w_0$ .

online training  $w_i \leftarrow w_i - \alpha \frac{\partial}{\partial w_i} \text{Loss}(w)$

- Then, plugging this back into Equation (1), and folding the 2 into the unspecified learning rate  $\alpha$ , we get the following learning rule for the weights:

-2 ถูกยึดไว้ให้  $\alpha$

$$w_0 = w_0 + \alpha (y - h_w(x))$$

$$w_1 = w_1 + \alpha (y - h_w(x)) \times x$$

Target value

Actual value



This is the weight updating for a single training example.


- For **N training examples**, we want to minimize the sum of the individual losses for every example (**batch gradient descent** learning rule).

batch training

ปัดขึ้นช้ากว่า

$$w_0 = w_0 + \alpha \sum_j (y_j - h_w(x_j))$$

$$w_1 = w_1 + \alpha \sum_j (y_j - h_w(x_j)) \times x_j$$



Sum derivatives of the loss values from N training examples.

# Multivariate linear regression

- We can easily extend to multivariate linear regression problems, in which each example  $x_j$  is an  $n$ -element vector. Our hypothesis space is:

$$h_{sw}(x_j) = w_0 + w_1 x_{j,1} + \dots + w_n x_{j,n} = w_0 + \sum_i w_i x_{j,i}$$

- The  $w_0$  term, *the intercept*, stands out as different from the others.
- If we introduce input attribute ( $x_{j,0}$ ) which is defined as always equal to 1, then  $h$  is simply the dot product of the weights and the input vector:

$$h_{sw}(x_j) = \mathbf{w} \cdot \mathbf{x}_j = \mathbf{w}^T \cdot \mathbf{x}_j = \sum_i w_i x_{j,i}$$

Diagram illustrating the dot product calculation:

Input vector  $\mathbf{x}_j = \begin{bmatrix} 1 \\ 3 \\ 4 \\ 1 \end{bmatrix}$  (where the first element is  $x_{j,0} = 1$ ) and weight vector  $\mathbf{w} = [1 \ 5 \ 2]$ .

Calculation:  $\begin{bmatrix} 1 \\ 3 \\ 4 \\ 1 \end{bmatrix} \cdot [1 \ 5 \ 2] \rightarrow [1 \ 3 \ 4 \ 1] \cdot [1 \ 5 \ 2] = 1 \cdot 1 + 3 \cdot 5 + 4 \cdot 2 + 1 \cdot 2 = 25$  (scalar)

Handwritten notes show the formula  $w_0 x_{j,0} + \sum_i w_i x_{j,i}$  and the calculation  $1 \cdot 1 + 3 \cdot 5 + 4 \cdot 2 + 1 \cdot 2 = 25$ .



- The update weight  $w_i$  by substitute  $h_{sw}(x_j)$  into the equation in page 7:

$$w_i = w_i + \alpha \sum_j x_{j,i} \times \left( \overset{\text{Target value}}{\downarrow} y_j - \overset{\text{Actual value}}{\swarrow} h_{sw}(x_j) \right)$$

Loss (multivariate)


- This process is **equivalent** in finding the best vector of weights,  $\mathbf{w}^*$ , that minimizes squared-error loss over the examples:

$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_j (y_j - h_{sw}(x_j))^2$$

↓  
get param  $\mathbf{w}$  which minimize loss on all  $x_j$   
(argument)

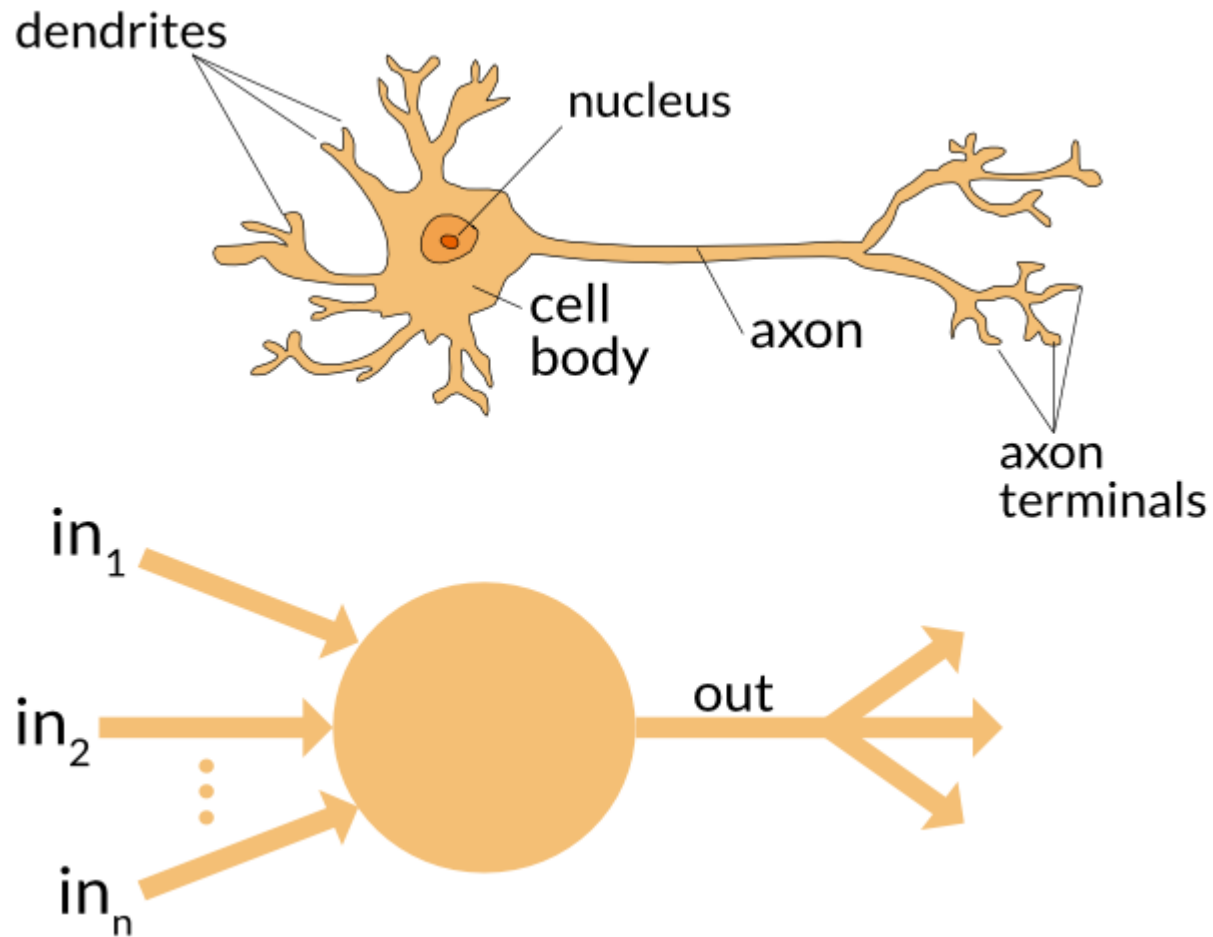
# Linear Classification With Perceptron

# Biological neuron




- Dendrites receive signals
- Axon sends signals out to other neurons

# Artificial neuron



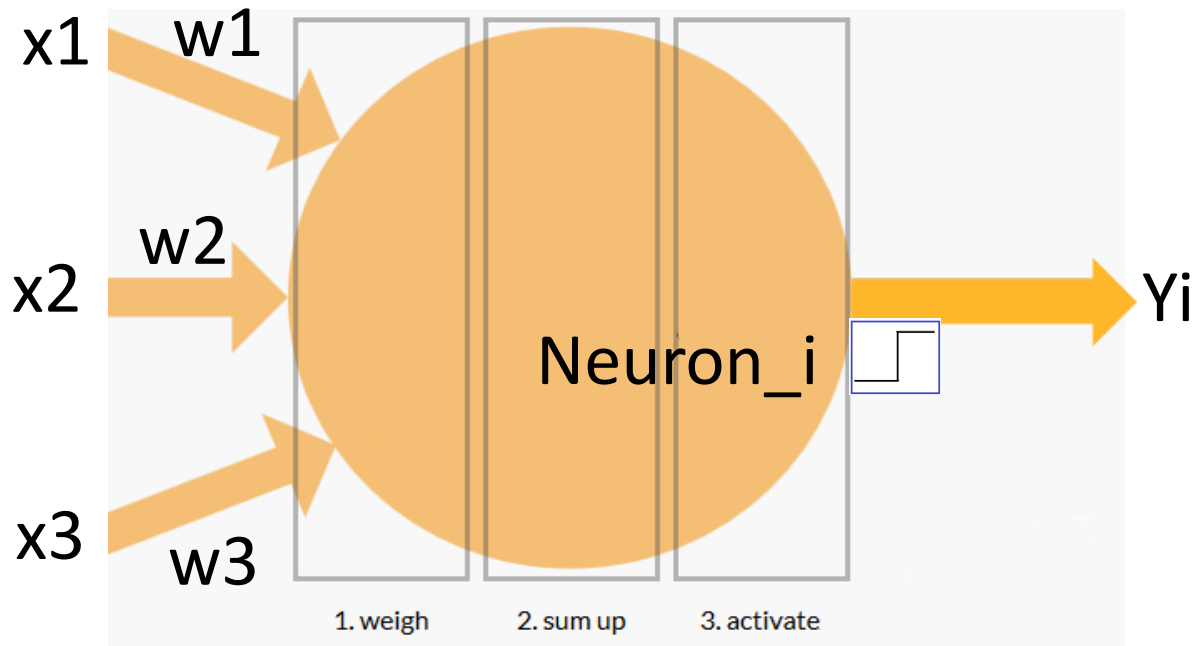
# Inside artificial neuron



We will call this artificial neuron as a **perceptron**.

# Notations

**Perceptron** is the simplest form of a neural network. It consists of a single neuron with adjustable weights and a hard limit activation function.



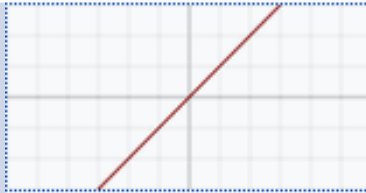
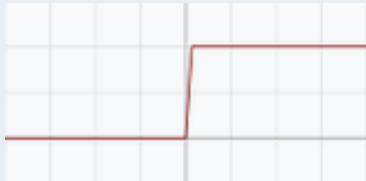
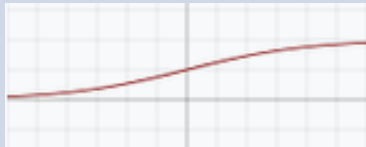
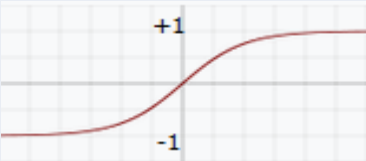
There are many kinds of **activation function**.



**Frank Rosenblatt**



invented perceptron in 1957

# Activation functions

Name	Equation	Plot
Identity (Linear)	$f(x) = x$	
Binary step	$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	
Sigmoid (Logistic)	$f(x) = \frac{1}{1 + e^{-x}}$	
TanH	$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	
Softmax	$f_i(\vec{x}) = \frac{e^{x_i}}{\sum_{j=1}^J e^{x_j}} \quad \text{for } i = 1, \dots, J$	

# Linear separability

Data in the n-dimensional space is **linearly separable** if two classes of data are divided by a hyperplane into two decision regions.



In general, the hyperplane is defined by the linearly separable function.

$$\sum_{i=1}^n x_i w_i - \theta = 0$$

(The threshold  $\theta$  can be used to shift the decision boundary.)



- The perceptron learns its classification task by making small adjustments in the weights to reduce the difference between the actual and desired (target) outputs of the perceptron.

$$e(p) = Y_d(p) - Y(p)$$

Target value    Actual value  
↓                      ↓

ข้อบกพร่อง

Where

$p$  is the training pattern (example)

$Y_d(p)$  is the desired (target) output of pattern  $p$


$Y(p)$  is the actual output

$e(p)$  is the difference (error) between  $Y_d(p)$  and  $Y(p)$

- It uses  $e(p)$  to update weights of the next iteration,

$$w_i(p + 1) = w_i(p) - \alpha \cdot x_i(p) \cdot e(p)$$

where  $\alpha$  is the learning rate (0~1)



# Perceptron learning algorithm

## Step 1: Initialization

- Set initial weights  $w_1, w_2, \dots, w_n$  and thresholds ( $\theta$ ) to small random numbers in the range  $[-0.5, +0.5]$
- $p = 1$  # the first pattern

## Step 2: Activation

$$Y(p) = \text{step} \left( \sum_{i=1}^n x_i(p) \cdot w_i(p) - \theta \right)$$

where  $n$  is the number of perceptron inputs.

## Step 3: Weight training by using gradient descent

$$w_i(p+1) = w_i(p) + \Delta w_i(p),$$

$$\Delta w_i(p) = -\alpha \cdot x_i(p) \cdot e(p)$$

$$\text{where } e(p) = Y_d(p) - Y(p)$$

$$w_1 = w_1 + \alpha (y - h_w(x)) \times x$$

Target value    Actual value

## Step 4: Iteration

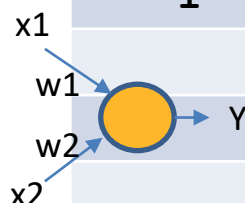
Increase  $p$  by one, go back to step 2 until each pattern is trained.

**Step 5:** If the perceptron doesn't converge,  $p = 1$  and repeat steps 2 – 4.

# Example: Train a perceptron on AND

$$w_i(p+1) = w_i(p) - \alpha \cdot x_i(p) \cdot e(p)$$

$\theta = 0.3, \alpha = 0.1$






Epoch	Inputs x1 x2	Yd (desired)	Weights w1 w2	Y (actual)	Error	Weights w1 w2
1	0 0	0	0.3 -0.1	0	0	0.3 -0.1
	0 1	0	0.3 -0.1	0	0	0.3 -0.1
	1 0	0	0.3 -0.1	1	-1	0.4 -0.1
	1 1	1	0.4 -0.1	1	0	0.4 -0.1
2	0 0	0	0.4 -0.1	0	0	0.4 -0.1
	0 1	0	0.4 -0.1	0	0	0.4 -0.1
	1 0	0	0.4 -0.1	1	-1	0.3 -0.1
	1 1	1	0.3 -0.1	0	1	0.2 -0.2
.	.	.	.	.	.	.
.	.	.	.	.	.	.
.	.	.	.	.	.	.
5	0 0	0	0.15 0.15	0	0	0.15 0.15
	0 1	0	0.15 0.15	0	0	0.15 0.15
	1 0	0	0.15 0.15	0	0	0.15 0.15
	1 1	1	0.15 0.15	1	0	0.15 0.15

$\Delta w_1 = -\alpha \cdot x_1(p) \cdot e(p) = 0.1$   
 $\Delta w_2 = -0.1 \cdot 0 \cdot (-1) = 0$

# Problem of perceptron



- Perceptron can learn only simple linear separable problems, e.g., AND, OR. It failed on XOR problem.



- A single perceptron can classify only linear separable problems, regardless of whether we use a hard-limit or soft-limit activation functions.
- Moreover, increasing the number of perceptrons in the same layer doesn't help.

# Linear Classification by Support Vector Machine

- SVMs (Vapnik, 1990's) choose the linear separator with the **largest margin**



V. Vapnik

- Good according to intuition, theory, practice
- SVM became famous when, using images as input, it gave accuracy comparable to neural-network


# Linear Classification by Support Vector Machine

$\exists \mathbf{w}$  such that  $\forall t$


$$y_t(\mathbf{w} \cdot \mathbf{x}_t) \geq \gamma > 0$$



Called the *margin*



# Linear Classification by Support Vector Machine



Suppose that the value of margin is 1.

We are asked to find a set of weights ( $w$ ) such that, for all pattern  $t$ ,

$$\text{for } y_t = +1, \quad w \cdot x_t + b \geq 1$$

$$\text{and for } y_t = -1, \quad w \cdot x_t + b \leq -1$$

That is, we want to satisfy all of the **linear** constraints

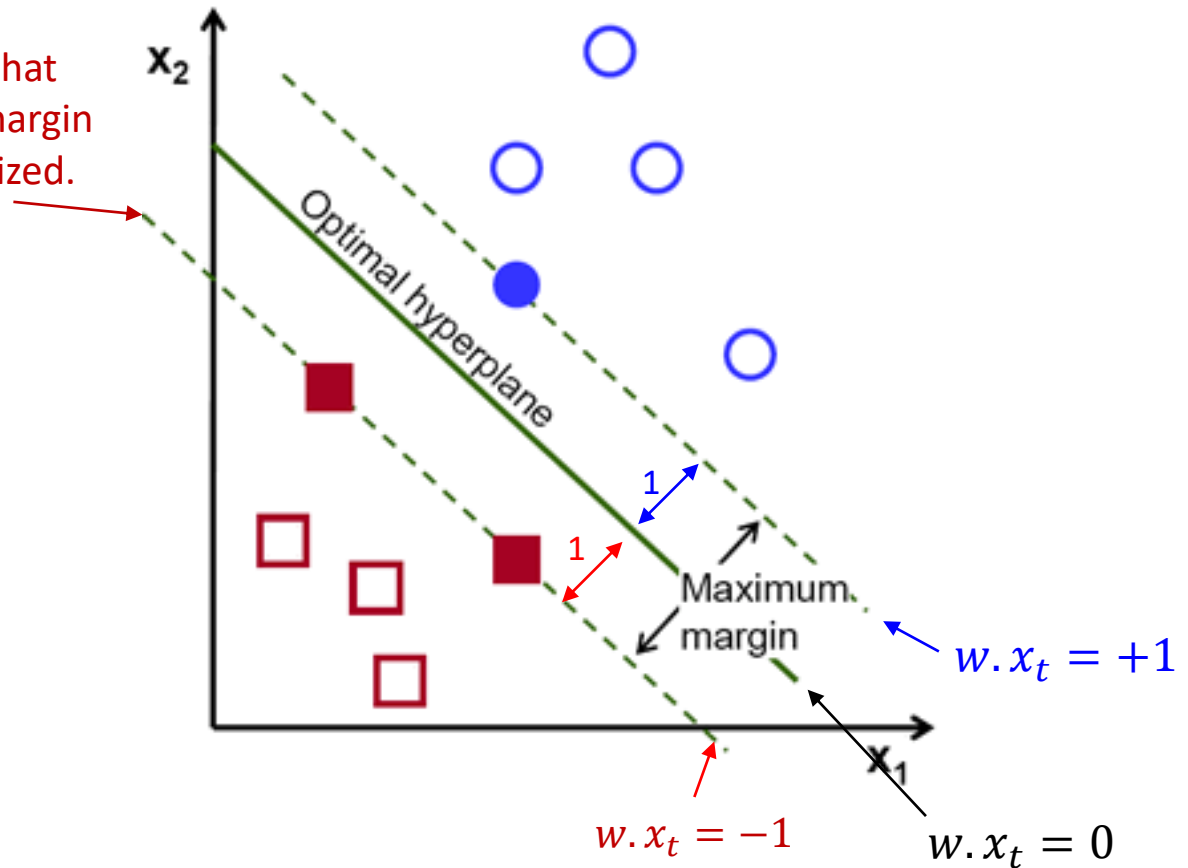
$$y_t (w \cdot x_t + b) \geq 1 \quad \forall t$$



# SVM Loss

- Loss of SVM


Any rectangle instance that crosses the maximum margin (dash line) will be penalized.



# SVM Loss

Penalty = 0

Penalty > 0



# SVM Loss

Let  $s = f(x_i, w)$  ; score of input  $x_i$   
 $s_{y_i}$  is the score of the target class



the SVM loss has the form:

Score of the target class  $s_{y_i}$  (green arrow)  
Score of other class  $s_j$  (red arrow)

$$L_i = \sum_{j \neq y_i} \begin{cases} 0 & \text{if } s_{y_i} \geq s_j + 1 \\ s_j - s_{y_i} + 1 & \text{otherwise} \end{cases}$$
$$= \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

cat	$s_{y_i}$ <b>3.2</b>
car	$s_j$ 5.1
frog	$s_j$ -1.7

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

$$\begin{aligned} &= \max(0, 5.1 - 3.2 + 1) \\ &\quad + \max(0, -1.7 - 3.2 + 1) \\ &= \max(0, 2.9) + \max(0, -3.9) \\ &= 2.9 + 0 \\ &= 2.9 \end{aligned}$$

- Next class
  - Multilayer neural networks