

Deep Learning in Natural Language Processing

Kietikul Jearanaitanakij

Department of Computer Engineering, KMITL

Representing Word with Numeric Values

1. One-hot encoding

Let n be the number of words in dictionary.

We use the n -dimension one-hot vector to represent each word.

Example:

Dict = { apple, orange, mango, durian }

apple = $\langle 1, 0, 0, 0 \rangle$

orange = $\langle 0, 1, 0, 0 \rangle$

mango = $\langle 0, 0, 1, 0 \rangle$

durian = $\langle 0, 0, 0, 1 \rangle$

Representing Word with Numeric Values

2. TF-IDF encoding

TF – Term frequency

$$\frac{\text{No. of times the term occurs in the document}}{\text{Total No. of terms in the document}}$$

Example: sentence1 = < I like orange >
sentence2 = < I like banana >
sentence3 = < I don't know fresh figs >

Let sentence = document, word = term.

In sentence1, $TF(I) = TF(like) = TF(orange) = 1/3 = 0.33$

In sentence2, $TF(I) = TF(like) = TF(banana) = 1/3 = 0.33$

In sentence3, $TF(I) = TF(don't) = TF(know) = TF(fresh) = TF(figs) = 1/5 = 0.2$

$IDF(I) = \log(3/3) = 0$; So, 'I' is not important term. $IDF(like) = \log(3/2) = 0.17$

$IDF(orange) = \log(3/1) = \log(3) = 0.47$; Important term!

$= IDF(banana), IDF(don't), IDF(know), IDF(fresh), IDF(figs)$

IDF – Inverse Document Frequency

$$\log\left(\frac{\text{Total No. of documents}}{\text{No. of documents that a term occurs}}\right)$$



Representing Word with Numeric Values

TF-IDF encoding (cont.)

sentence1 = < I like orange > , TF of each word = 0.33

sentence2 = < I like banana > , TF of each word = 0.33

sentence3 = < I don't know Figs and Plum> , TF of each word = 0.2

Note: *TF of each word in a sentence is identical by incidental.*

	banana	don't	figs	fresh	know	like	orange
IDF	0.47	0.47	0.47	0.47	0.47	0.17	0.47

Note: *IDF of each word has only a single value. ('I' was removed since its IDF=0)*

Word positions in vector (sorted): [banana, don't, figs, fresh, know, like, orange]

Word vector for each sentence is given below:

Sentence1: [0, 0, 0, 0, 0, 0.33*0.17, 0.33*0.47]

Sentence2: [0.33*0.47, 0, 0, 0, 0, 0.33*0.17, 0]

Sentence3: [0, 0.17*0.47, 0.17*0.47, 0.17*0.47, 0.17*0.47, 0, 0]

Warning! Result from Scikit-learn may different because it applies some advanced techniques, e.g. smoothing, log base e.

Representing Word with Numeric Values

3. Word Vector

- The most popular word embedding.
- Unlike one-hot and tf-idf encodings, word vector can capture the meaning, semantic similarity, and relationship of a word and its surrounding words.
- Each word is encoded as a multi-dimensional vector.

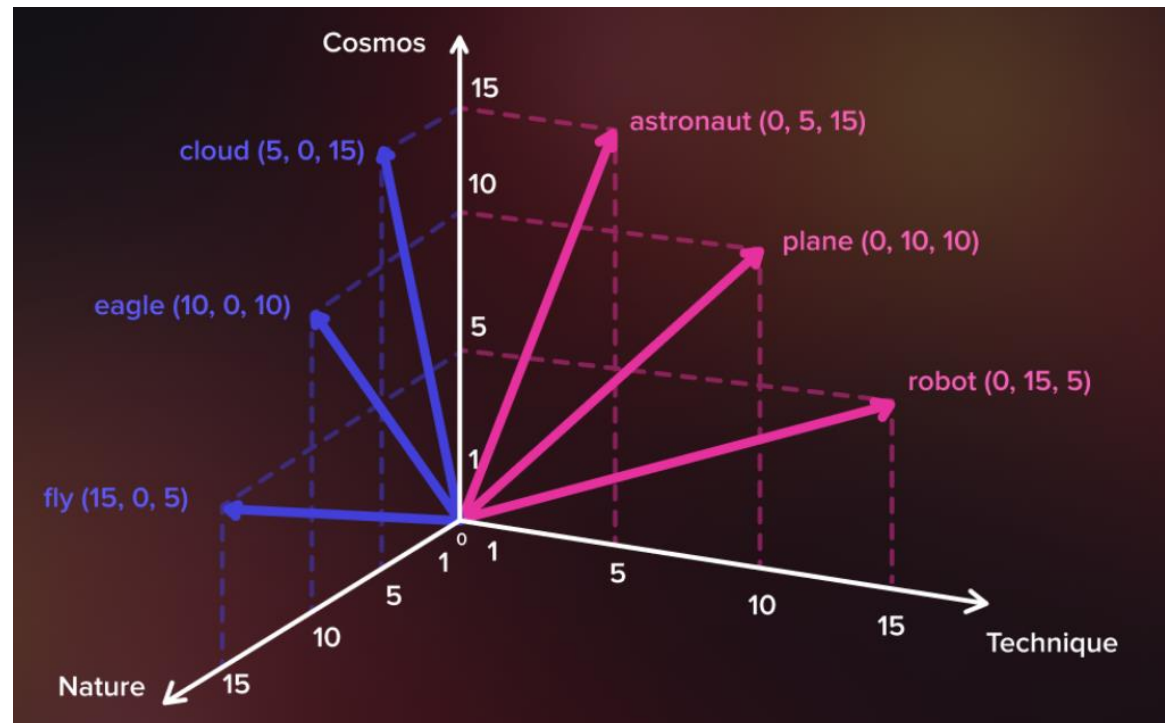


Figure credit: <https://serokell.io/blog/word2vec>

Representing Word with Numeric Values

- Word2Vec uses a **shallow neural network** (have one or two hidden layers between the input and output) to learn the meaning of words from a large corpus of texts.
- The shallow neural network of Word2Vec can quickly recognize semantic similarities and identify synonymous words using **logistic regression classifier**.
- To represent a particular word as an input vector in multidimensional space, simply use a bag-of-word encoding.
- The training algorithm is categorized into two modes: **continuous bag of words (CBOW)** or **skip-gram**.

Representing Word with Numeric Values

- **Continuous bag of words (CBOW)** => **Tom** **?** **Goong**

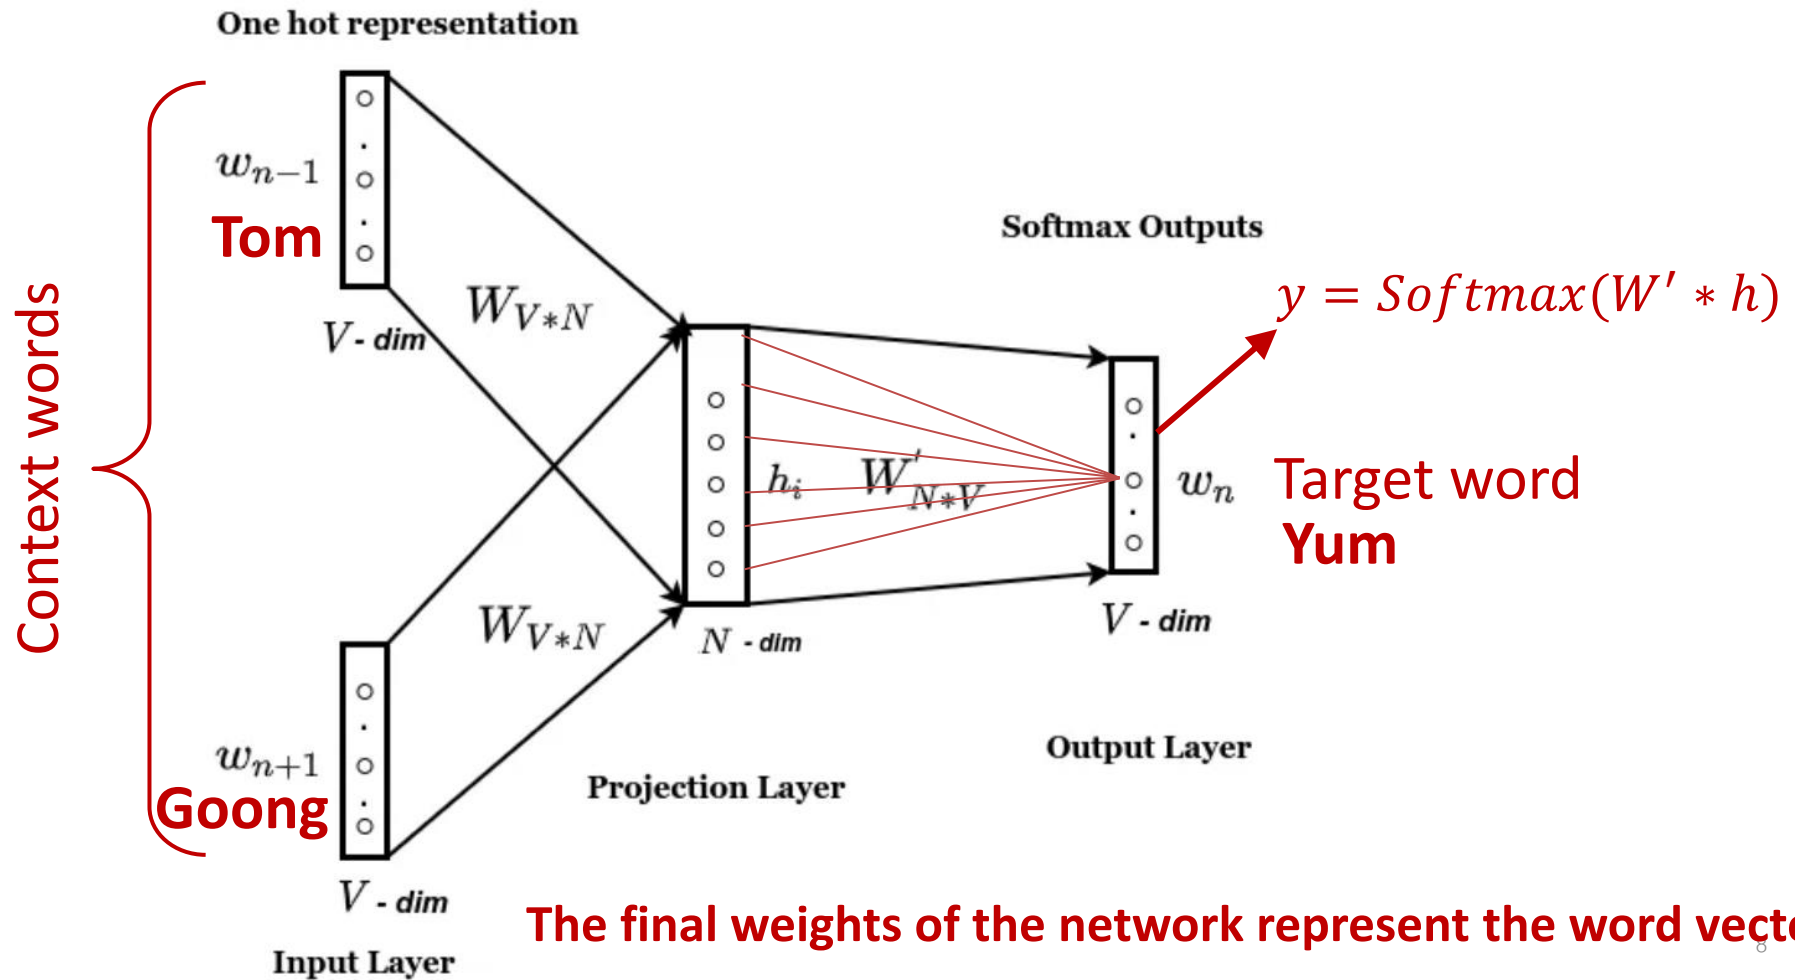
Predict the target word 'Yum' from the context words (Tom, Goong)

- The **input layer** is represented by a **one-hot encoded vector**, where each element in the vector corresponds to a specific word in the vocabulary. For example, if the vocabulary contains 10,000 words, the input layer will have 10,000 elements.
- The **hidden layer** is where the word embeddings are learned. The number of neurons in the hidden layer can be varied depending upon the dimension of the word vector.
- The **output layer** is also a dense layer, with each neuron representing a specific word in the vocabulary. The number of neurons in the output layer is the same as the number of words in the vocabulary.

Representing Word with Numeric Values

- Continuous bag of words (CBOW)

Tom ? Goong




Representing Word with Numeric Values

- Skip-gram => ? **Yum** ?.

Predict the context words ('Tom', 'Goong') from the given word 'Yum'.

One-hot encoding
(Dimension = Vocab size)



No. of neurons = dimensions of word vector

Representing Word with Numeric Values

According to the original paper, Mikolov et al., it is found that

- **Skip-Gram** works well with small datasets, and can better represent less frequent words (e.g. pulchritudinous).
- **CBOW** is found to train faster than Skip-Gram, and can better represent more frequent words (e.g. beautiful, nice).

Representing Word with Numeric Values

Interesting facts about Word2Vec

- Word vectors are meaningful as you can find their relationships.
$$\text{Vector}(\text{King}) - \text{Vector}(\text{Man}) + \text{Vector}(\text{Woman}) = \text{Vector}(\text{Queen})$$
- **Gensim**¹ provides an implementation of **Word2Vec** that you can train the model from scratch.
- PyThaiNLP² provides **Thai2Vec** which contains 51,556 Thai word embeddings each of 300 dimensions.
- Another well-known word encoding is **GloVe** (Global Vectors for Word Representation) from Stanford University which was trained on different method but give similar results.

1. <https://radimrehurek.com/gensim/models/word2vec.html>

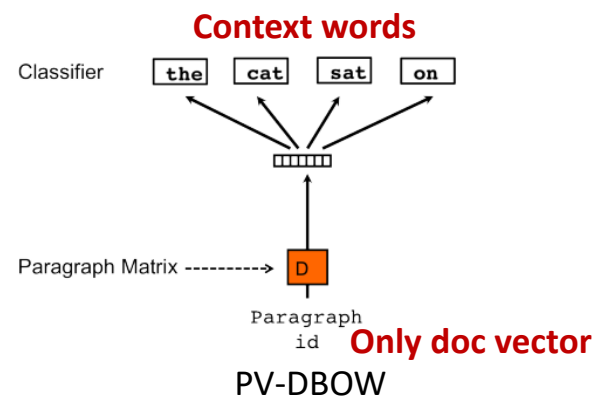
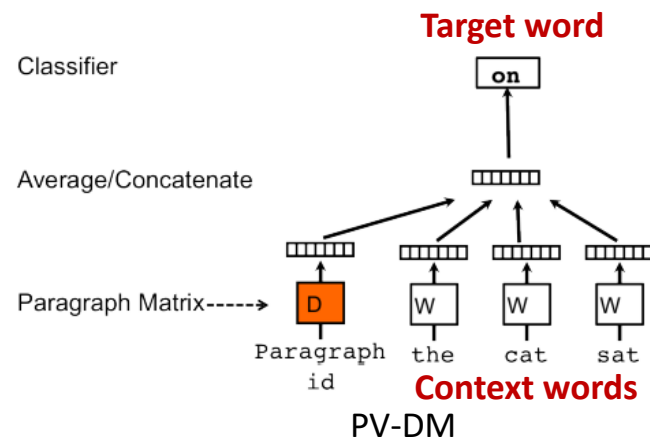
2. https://pythainlp.github.io/tutorials/notebooks/word2vec_examples.html

Representing Sentence with Numeric Values

- The simple strategy to find the sentence vector is to **average all vectors of words** in that sentence.
- Although this simple averaging works in most task, it performs poorly for sentiment analysis tasks, because it "loses the word order in the same way as the standard bag-of-words models do".
- Other advanced methods to find a sentence vector can be found in:
 - https://nlp.stanford.edu/~socherr/EMNLP2013_RNTN.pdf
 - <https://arxiv.org/abs/1506.06726>
 - <https://aclanthology.org/P16-1089.pdf>
 - <https://arxiv.org/pdf/1908.10084.pdf>

Representing Document with Numeric Values

- **Doc2Vec** has a similar training process to Word2Vec. Two Doc2Vec training algorithms: PV-DM and PV-DBOW. Unseen document is inferred its vector by the average of word vectors.
- **Distributed memory (PV-DM)** : instead of using just words to predict the target word, we also added another feature vector, which is document-unique (paragraph id). The target word is predicted. (PV-DM consumes more memory, but has better result than PV-DBOW)
- **Distributed bag of words (PV-DBOW)** : Only the document vector (or paragraph id) is fed into the hidden layer. The context words in the document are predicted. (PV-DBOW is faster than PV-DM and works very well on shortish-docs)
- Once words are represented as vectors of numeric values, we can apply them to any deep learning models.



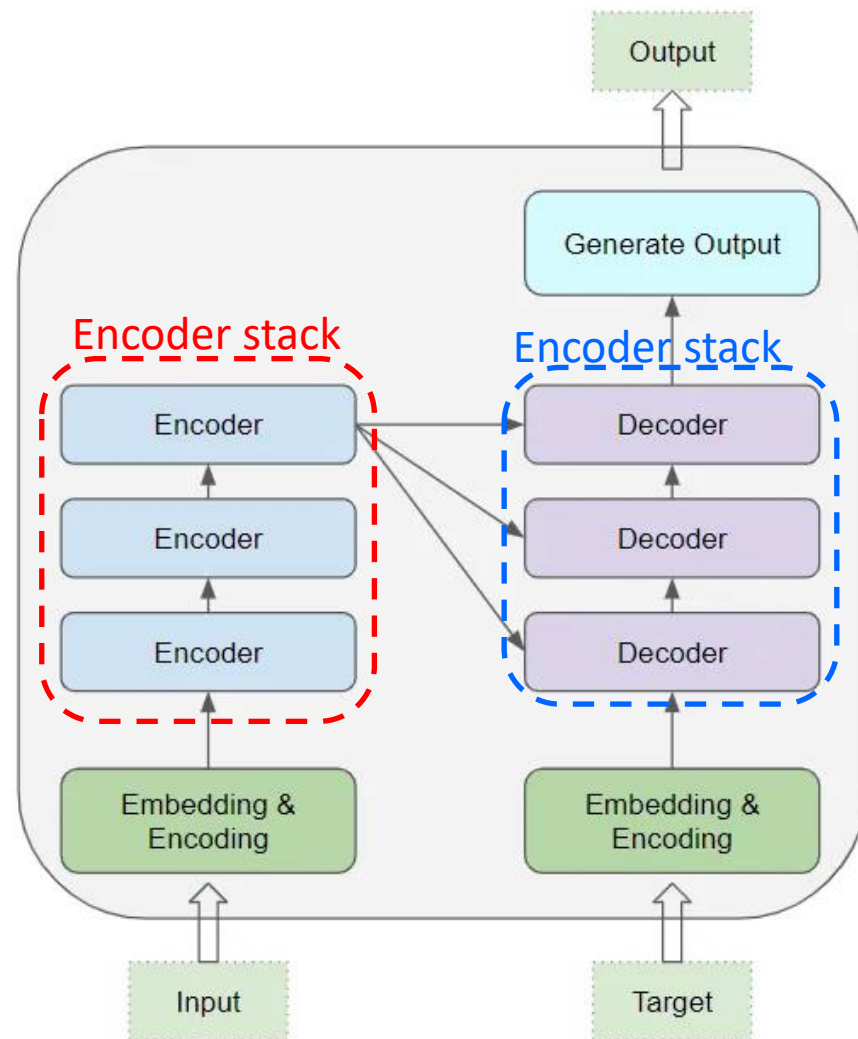
Transformer

- **Transformer** is a deep learning architecture that uses **Attention Mechanism** to significantly improve the performance of NLP translation models.
- It was first introduced in the paper “**Attention is all you need**”¹
- Numerous projects including **Google’s BERT** and **OpenAI’s GPT** series have built on this foundation
- Transformer takes a text sequence as input and produce another text sequence as output.



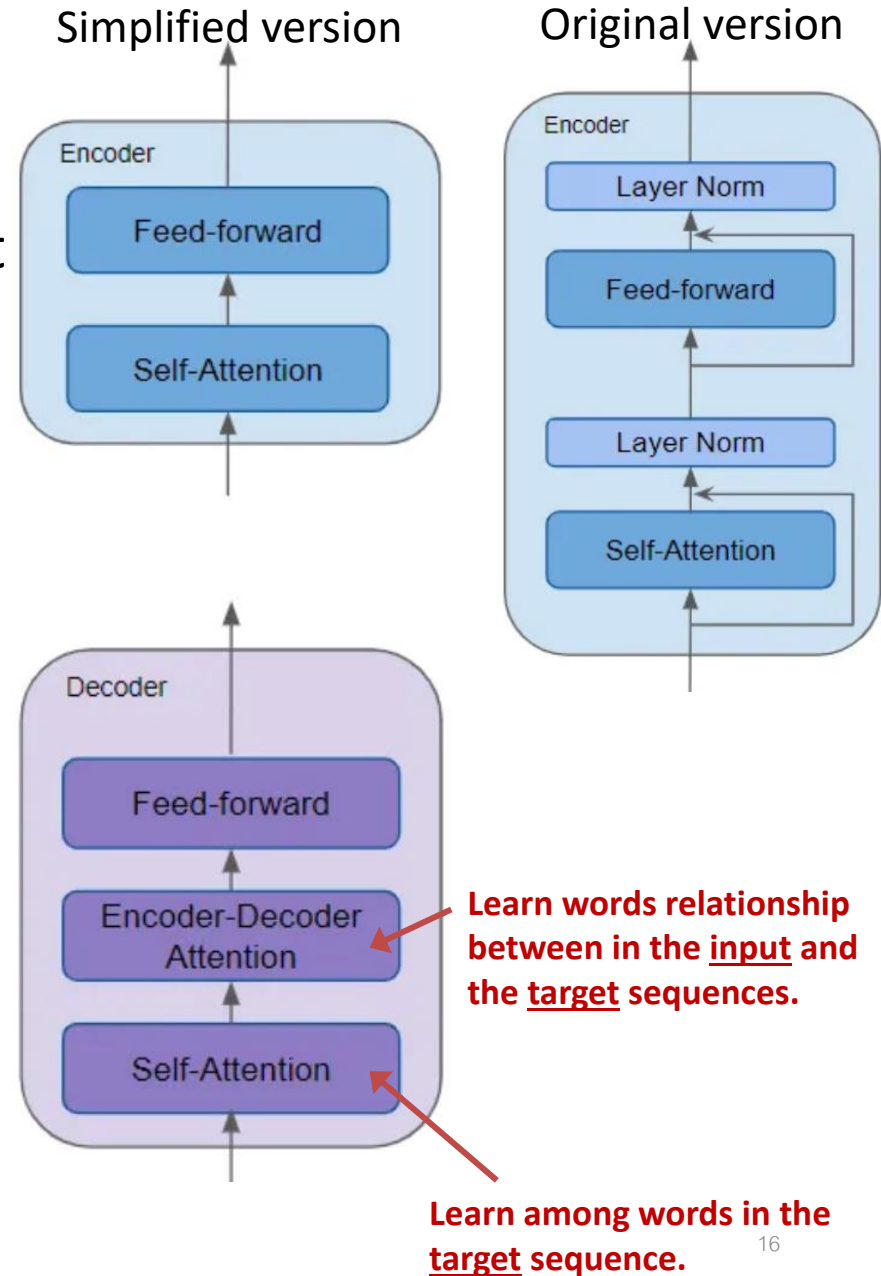
Transformer

- It contains a stack of **Encoder layers** and **Decoder layers**.
- Each layer in Encoder and Decoder layers has its own set of weights.
- These weights can learn relationship among each word and other words in the sentence.
- Finally, there is an Output layer to generate the final output.



Transformer

- The **Encoder** contains the all-important **Self-attention layer** that computes the relationship between different words in the sequence, as well as a Feed-forward layer.
- The **Decoder** contains the Self-attention layer and the Feed-forward layer, as well as a second Encoder-Decoder attention layer.



Transformer

- The most important function in Transformer is its **Attention Mechanism**.
- While processing a word, Attention enables the model to focus on other words in the input that are closely related to that word.

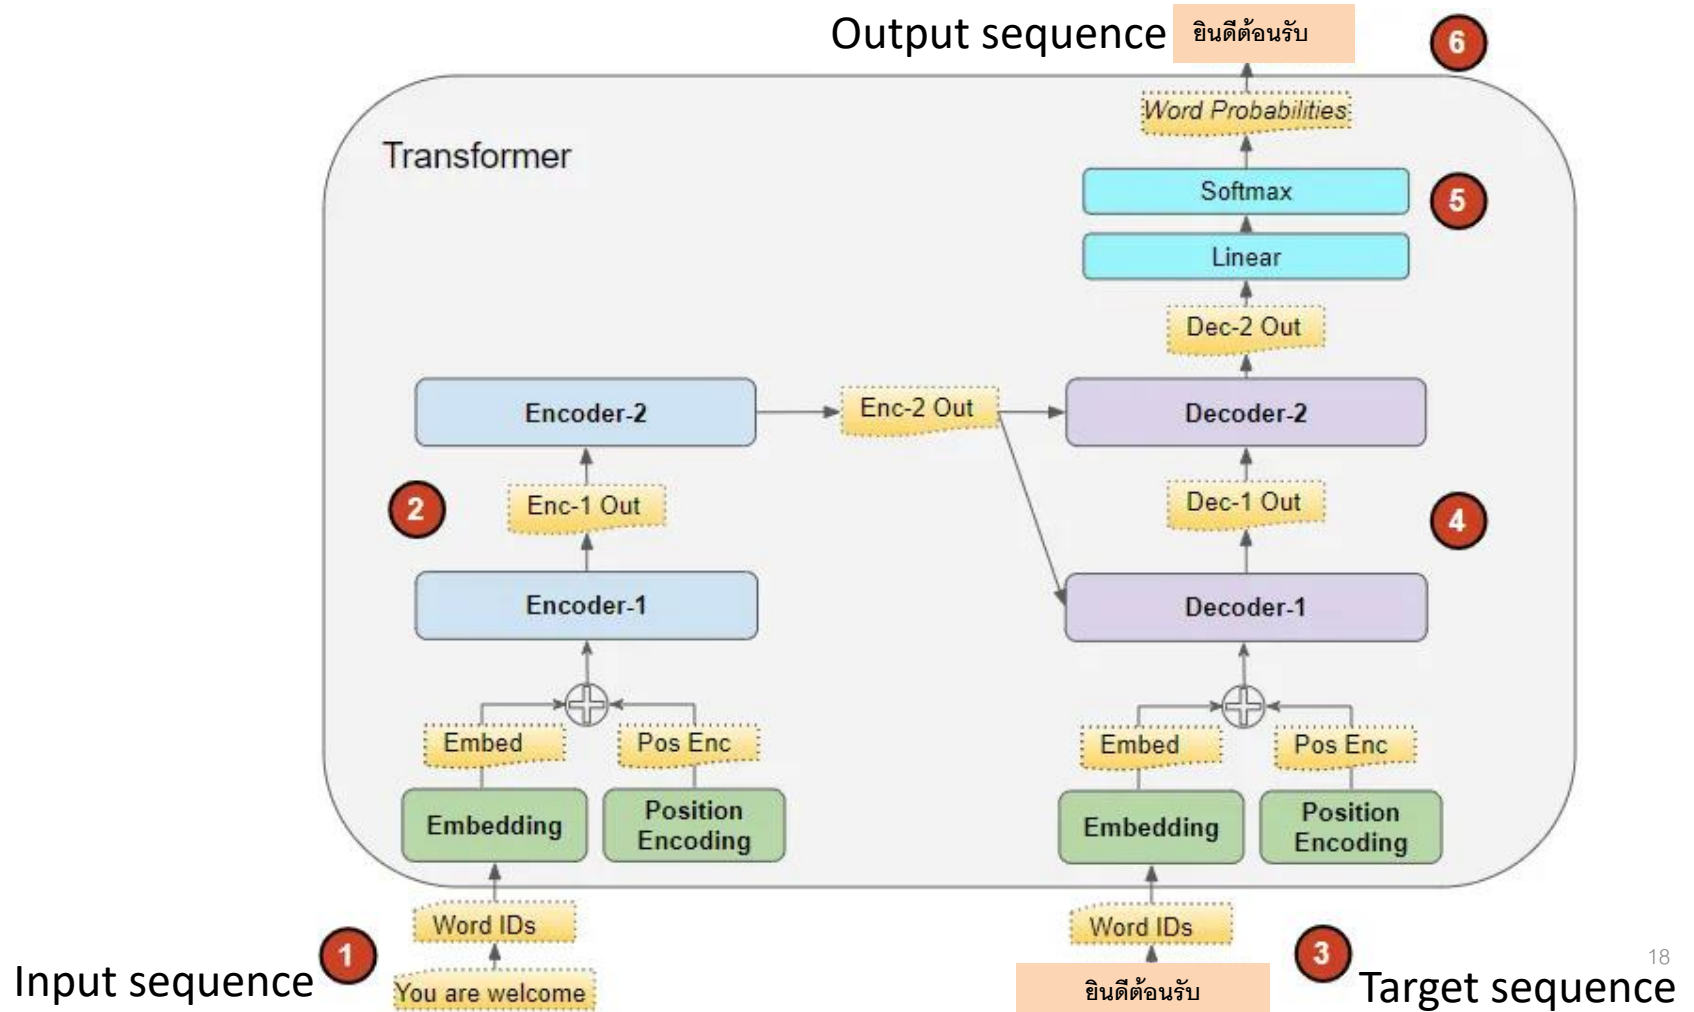


'Ball' is closely related to **'blue'** and **'holding'**. On the other hand, **'blue'** is not related to **'boy'**

- Transformers include multiple attention scores for each word.
- Training data consists of two parts:
 1. The source or input sequence (e.g. "You are welcome" in English, for a translation problem)
 2. The destination or target sequence (e.g. "ยินดีต้อนรับ" in Thai)

Transformer

- The Transformer's goal is to learn how to output the target sequence, by using both the input and target sequences.



Transformer

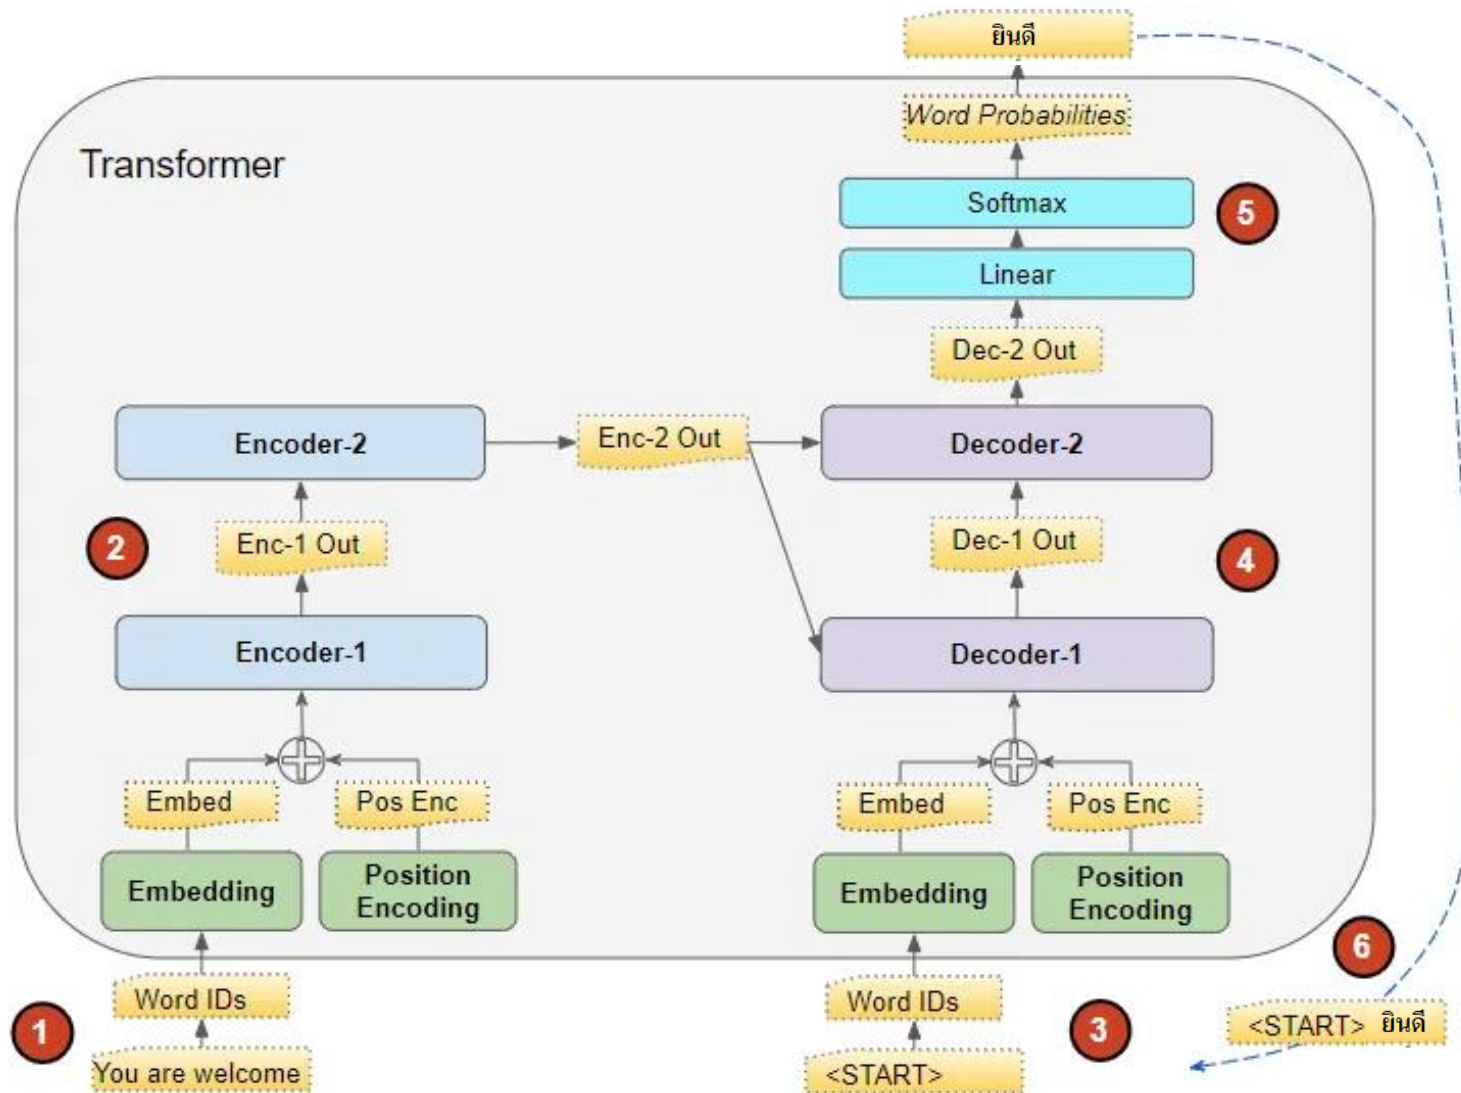
Training steps in Transformer.

1. The **input sequence** is **encoded with the position of each word**. Then the encoded signal is fed to the Encoder.
2. **Encoder stack** processes this signal with **self attention**.
3. The **target sequence** is **prepended** with a start-of-sentence token, converted into target embeddings (with positional encoding), and fed to the Decoder.
4. The **stack of Decoders** processes the target sequence embedding with **self attention**.
5. The outputs from the **Decoder stack** and the **Encoder stack** are brought into the **attention** calculation to find the relationship between them. (Encoder-Decoder attention)
6. The Output layer converts it into word probabilities (softmax function) and the final **output sequence**.
7. Loss value is calculated and the gradient is back-propagated.

Inference in Transformer

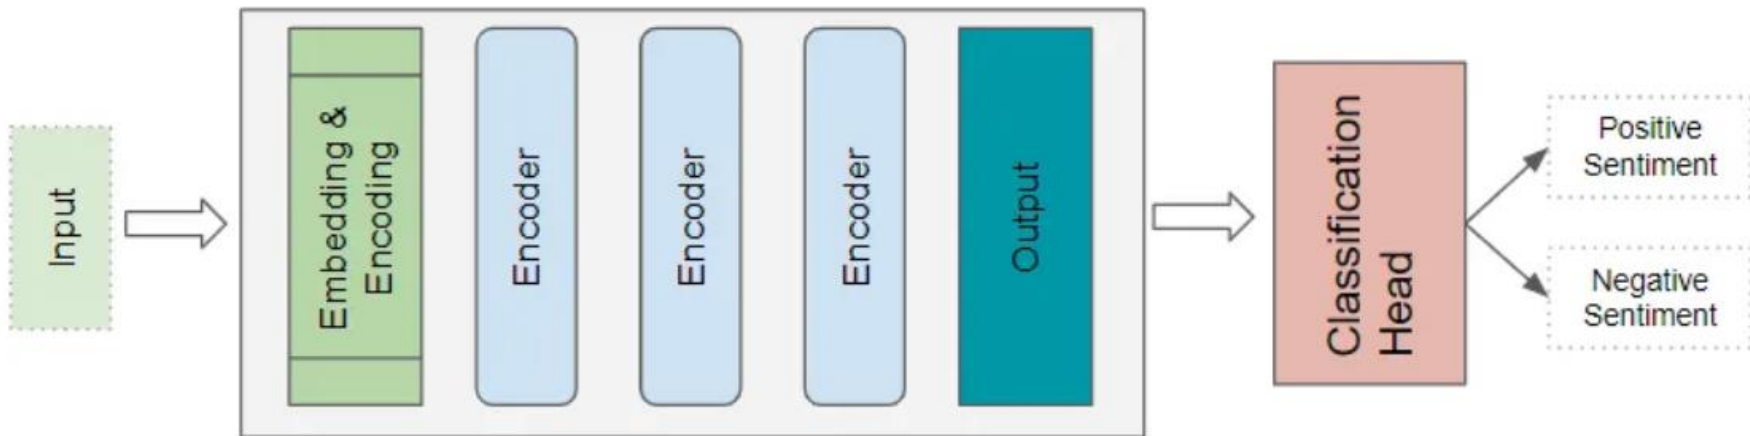
- During **Inference**, we have only the **input sequence** and **don't have the target sequence** to pass as input to the Decoder.
- The goal of the Transformer is to produce the target sequence (at the output layer) from the input sequence alone.
- We **generate the output in a loop** : feed the **entire** output sequence from the previous timestep to the Decoder in the next timestep until we reach an end-of-sentence token <eos> at the output layer.

Inference in Transformer



Applications of Transformer

- Transformer is frequently used in sequence-to-sequence models for applications such as **Machine Translation**, **Text Summarization**, **Question-Answering**, **Named Entity Recognition**, and **Speech Recognition**.
- The basic Encoder Layer is used as a common building block with different application-specific '**heads**' depending on the problem being solved.



Sample application of transformer: Sentiment analysis

Large Language Model (LLM)

- **LLM** is a very large deep learning models that are pre-trained on vast amounts of data.
- LLM can perform completely different tasks such as **answering questions**, **summarizing documents**, **translating languages**, **sentiment analysis**, and **text generation**, etc.
- Once trained, LLMs can be finetuned to perform specific task by using small data.

Three common learning models exist:

- **Zero-shot learning**; Base LLMs can respond to a broad range of requests without explicit training or prompts.
Ex: Query “What is the capital of France?”, Answer: “Paris”
- **Few-shot learning**: By providing a few relevant prompt examples, base model performance significantly improves in that specific area.
Ex: “France: Paris, Germany: Berlin, Canada: Ottawa”, Query: “Thailand”
- **Fine-tuning**: This is an extension of few-shot learning in that we train a base model with a specific task’s dataset.